



[Powered by Prosus and AISO]

[May 18th, 2026]

Hackathon Cheatsheet

Below you can find answers to most of your questions regarding the hackathon, including logistics, schedule, submissions, judging criteria, and available technical resources.

***PLEASE NOTE – Google Docs often do not render properly on phones!
Please review this document on your laptop.***

Hackathon overview

This is an agentic AI hackathon focused on autonomous business operators. Your challenge is to build an AI agent capable of independently managing and optimizing a simulated restaurant environment.

Your agent will make decisions around pricing, staffing, inventory, supplier management, customer satisfaction, and overall business strategy, with the goal of maximizing both profit and operational performance over a long-term simulation.

The environment is dynamic and unpredictable, meaning there is no hardcoded solution. The strongest agents will be adaptive, strategic, and capable of making autonomous decisions across changing scenarios.

What counts as an AI agent? We are not looking for simple chat-based co-pilots or scripted workflows. We are looking for autonomous systems that can reason, act, and operate on behalf of the user.

Teams will be evaluated primarily based on the performance of their agents in a hidden evaluation environment. The top teams will advance to final presentations, where judges will also assess technical understanding, creativity, and implementation quality.

The team size is between 3–5 people. Good luck, and may the best autonomous operator win 🚀

Repository

What free technical resources are available?

We will provide you with free technical resources to help you build your AI agent project. This includes access to a virtual machine and the GPT-4 and GPT-5 series of models. You are not required to use any of these resources! Here's how to access these resources:

- If you would like to use the models, you can request it via the [Discord channel](#) at `#request-your-api-key`.

What are the prizes for the winning teams?

The top three teams will receive a fast-tracked recruitment opportunity into Prosus and a €3,000 prize pool split into...

First: €1,500

Second: €1,000

Third: €500

What is the schedule for the hackathon?

7:30–8:00	Check-in & coffee	Network in the lounge, pick up badge
8:00–8:45	Opening ceremony	AISO & Prosus presentation + challenge intro
8:45–9:00	Team formation	Solo participants matched into teams (max 5)
9:00–17:00	Hacking	
12:00–13:00	Lunch	Food provided
17:00	Submission deadline	Tools down
17:00–18:00	Recruiter booth	Meet Prosus recruiters
17:00–18:00	Dinner + judging	Judges review submissions, select top 5 finalists
18:00–18:15	Top 5 teams reveal	Top 5 finalists are announced
18:30–19:15	Final presentations	Top 5 teams pitch to the full room
19:15–19:30	Winner selection	Judges deliberate — 1st, 2nd, 3rd
19:30–19:45	Closing & winner announcement	Results on stage
19:45–21:15	Networking drinks	Celebrate with the community

PLEASE KEEP YOUR ID CARD AND VISITOR PASS ON YOU AT ALL TIMES!

How do I upload pictures?

Please submit any pictures you take to the public folder [via this link](#).

How do I access the Wifi?

The Wifi does not have a password. The wifi network is **AI-House Guest**.

How do I join the hackathon Discord channel?

You can join the Discord channel [via this link](#).

What is the number to contact building security?

You can contact them via 06 50822027

How do we define a good solution?

Judging happens in **two stages**, so a "good solution" needs to clear two bars.

Stage 1 — Profit gets you into the top 5

After submissions close at 17:00, judges select the top 5 teams based on **leaderboard standing across all 30 evaluation games** (10 scenarios × 3 seeds). No subjective judgment here — just your aggregate score. Concretely:

- **Survive every game.** Going bankrupt scores -100,000 and will tank your aggregate. The single most important rule.
- **Be consistent.** Final ranking averages across all 30 cells. An agent that scores +40k on `baseline` but bankrupts on a hidden scenario will lose to a steady +5k everywhere.
- **Adapt — don't memorise.** 6 of the 10 final scenarios are *hidden*. Your agent has never seen them. It needs to react to what's in `alerts` and observable drift in the data.
- **Complete the matrix.** Partial submissions rank below complete ones. All 30 games must finish.

Score landmarks:

- **do_nothing** baseline: -100,000 (bankrupt by ~day 14)
- **naive_rule** baseline: around -15,000 (survives, doesn't optimise)
- 0 or positive: you've beaten the naïve baseline
- +15,000 to +35,000: competitive

Stage 2 — How you built it decides the top 3

If you make the top 5, you'll pitch for **5 minutes + 3 minutes Q&A**. Judges score on four criteria:

Criterion	What judges are looking for
Agent autonomy	Does it think, plan, and act with minimal human input — or is it a thinly-wrapped if/else?
Impact	Real-world usefulness of the approach; relevance to the autonomous-agents theme
Technical quality	Clean architecture, robust under unseen scenarios, code that holds up
Creativity	Novel approach or unexpected use of the tools/APIs

What good agents tend to do (none of these are required — they're patterns that work):

- Read **dishes_unavailable_at** and **pending_orders** every turn. Most failures come from teams missing these two fields.
- Use **save_notes** as memory. LLMs have no memory between turns otherwise.
- **React to alerts**. Scenario events announce themselves there — free intelligence.
- **Mix rules and LLM**. Cheap rules for must-do safety (don't bankrupt, reorder when low); LLM for judgment calls (pricing, promotions, recovery strategy).
- **Test across all 4 known scenarios + multiple seeds**. An agent that aces baseline on seed 42 alone is fragile.

For the pitch, be ready to:

- Answer hard and technical questions about your solution
- Your architecture in one diagram
- One specific thing your agent does that you're proud of
- How it handled a hard scenario you didn't expect

How do I submit a project?

Note that if your team does not feel that your project is finished, and would like to pass on submitting the project for judging, that's fine! You are not required to submit a project. We know that sometimes you are not able to complete your project in time.

In order to submit, you must complete two steps:

- Have a valid submission in the [leaderboard](#)
- **Fill in [this form](#).**

The leaderboard auto-updates as your agent plays games. Whatever your aggregate score is at 17:00 is what judges use to pick the top 5.

Here are the key things you need to include in your submission:

- Your final repo must be public on Github, and the last commit must be before 17:00!
- Please provide a short summary of what you would envision as the next steps via the form. A short summary of what you would build next will help the judges understand your project and vision.

What to do during the day

1. **Pick a unique team name** when you create your first game and **use it consistently** across every game your agent plays. The leaderboard groups by team name.
2. **Develop on the 4 known scenarios** (baseline, supply_crisis, tourist_season, renovation) with dev seeds 42, 88, 123. Iterate until your agent is robust across all four.
3. **[Watch the live dashboard](#)** — it auto-refreshes every 15 seconds and shows where you stand.

When the final evaluation phase opens (around 16:00)

The organisers will:

- **Reset the leaderboard.** All dev-phase scores are wiped — only evaluation-phase games count.
- **Unlock 6 hidden scenarios.** You'll suddenly see 10 scenarios total.
- **Set the evaluation matrix: 10 scenarios × 3 seeds = 30 games per team.** Different seeds from dev phase.

Once the matrix opens, you need to run all 30 games before 17:00. Some logistics:

- **Max 5 concurrent games per team.** Plan your runs.
- **Best score per cell counts** — you can retry a (**scenario**, **seed**) cell and keep the better result.
- **Partial matrices rank below complete ones.** Finishing all 30 beats almost-finishing with one great score.
- **Don't keep iterating on prompts in this phase.** Lock your agent before the matrix opens. Use the time to run.

The `evaluate.py` helper in the starter kit can drive the full matrix for you (check README file for more info):

Shell

```
python -m agents.evaluate agents.my_agent --seeds 7,55,99
```

If you make the top 5

Top 5 finalists are announced at **18:00**. You'll pitch at **18:30**:

- **5 minutes** of pitch + **3 minutes** of Q&A
- Be ready to show your code on screen — judges will ask about architecture
- Bring a laptop with your agent runnable for a quick live demo if needed

Meet the team from AISO and the mentors/team from Prosus HQ!



Vitor Castro, Student at VU and part of AISO



Jan Pecka, Student at UvA and part of AISO



Filip Szturo, Student at UvA and part of AISO



Amina Akhmedova, Student at VU and part of AISO



Tatjana Obenaus, Head of Employer Brand, Prosus



Andreea Tache, Marketing & Communications, Prosus



Daniella Linera Roose, Exec. Talent Acquisition Coordination, Prosus



Jannica Heibel, Marketing & Communications, Prosus



Gabriel Inyang, Global Talent Partner, Prosus



Olga Sokolva, Talent Acquisition Partner, Prosus

Meet the mentors



Asad Ismail, Senior AI Engineer , Prosus



Yingying Deng, AI Talent Lab Resident, Prosus



Isha Agrawal, AI Engineer, Prosus



Gowtham Venkatesan, AI Talent Lab Resident, Prosus



Nikolas Stavrou, AI Talent Lab Resident, Prosus



Marin Marian, AI Talent Lab Resident, Prosus



Saunaq Chakrabarty, AI Talent Lab Resident, Prosus